

Chapter 3 & Appendix B (continued)

Arithmetic for Computers

Floating Point

- Representation for non-integral numbers
 - Including very small and very large numbers
- Like scientific notation
 - -2.34×10^{56} ← normalized
 - $+0.002 \times 10^{-4}$ ← not normalized
 - $+987.02 \times 10^9$ ← not normalized
- In binary
 - $\pm 1.xxxxxxx_2 \times 2^{yyyy}$
- Types `float` and `double` in C

Floating Point Standard

- Defined by IEEE standard 754-1985
- Developed in response to divergence of representations
 - Portability issues for scientific code
- Now almost universally adopted
- Two representations
 - Single precision (32-bit)
 - Double precision (64-bit)

IEEE Floating-Point Format

single: 8 bits
double: 11 bits

single: 23 bits
double: 52 bits

| | | |
|---|----------|----------|
| S | Exponent | Fraction |
|---|----------|----------|

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0 \Rightarrow non-negative, 1 \Rightarrow negative)
- Normalized significand: $1.0 \leq |\text{significand}| < 2.0$
 - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
 - Significand is Fraction with the "1." restored
- Exponent: actual exponent + Bias
 - Ensures exponent is unsigned
 - Single: Bias = 127; Double: Bias = 1023

Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
 - Exponent: 00000001
 \Rightarrow actual exponent = $1 - 127 = -126$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
 - exponent: 11111110
 \Rightarrow actual exponent = $254 - 127 = +127$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
 - Exponent: 000000000001
 \Rightarrow actual exponent = $1 - 1023 = -1022$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
 - Exponent: 11111111110
 \Rightarrow actual exponent = $2046 - 1023 = +1023$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

Floating-Point Precision

- Relative precision
 - all fraction bits are significant
 - Single: approx 2^{-23}
 - Equivalent to $23 \times \log_{10} 2 \approx 23 \times 0.3$
 ≈ 6 decimal digits of precision
 - Double: approx 2^{-52}
 - Equivalent to $52 \times \log_{10} 2 \approx 52 \times 0.3$
 ≈ 16 decimal digits of precision

Binary Refresher

- When we look at a number like 10110_2 , we're seeing it as:

$$1(2^4) + 0(2^3) + 1(2^2) + 1(2^1) + 0(2^0)$$

$$= 16 + 4 + 2 = 22_{10}$$

Binary Decimal Points

- In decimal, 12.63 is the same as

$$1(10^1) + 2(10^0) + 6(10^{-1}) + 3(10^{-2})$$

- In binary, 101.01_2 is the same as

$$1(2^2) + 0(2^1) + 1(2^0) + 0(2^{-1}) + 1(2^{-2})$$

$$= 4 + 1 + 0.25 = 5.25$$

Useful Exponent Identities

- $a^b * a^c = a^{b+c}$
 - Why memorize more than 2^{10} when we can just break them down?
 - $2^{35} = 2^5 * 2^{30}$
 $= 2^5 * 2^{10} * 2^{10} * 2^{10} = 32 \text{ GB}$
- $a^{-b} = \frac{1}{a^b}$
 - When we have decimal terms, instead of seeing 2^{-1} , 2^{-2} , etc. use $\frac{1}{2^1}$, $\frac{1}{2^2}$, etc.

More Binary

- Dividing (shifting right) by 2_{10} is the same as moving the decimal point one place to the left.
 - Same reasoning as when we divide by 10 in base 10
- Multiplying (shifting left) works the same way

Floating-Point Example

- Represent -0.75
 - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
 - $S = 1$
 - Fraction = $1000...00_2$
 - Exponent = $-1 + \text{Bias}$
 - Single: $-1 + 127 = 126 = 01111110_2$
 - Double: $-1 + 1023 = 1022 = 01111111110_2$
- Single: $1011111101000...00$
- Double: $1011111111101000...00$

Floating-Point Example

- What number is represented by the single-precision float

11000000101000...00

- $S = 1$
- Fraction = 01000...00₂
- Exponent = 10000001₂ = 129
- $$\begin{aligned}x &= (-1)^1 \times (1 + 0.01_2) \times 2^{(129 - 127)} \\&= (-1) \times 1.25 \times 2^2 \\&= -5.0\end{aligned}$$

IEEE 754-1985 Specials

- We reserve all 0s and all 1s in the exponent. This is why:
- $0111111110000...00 = +\infty$
- $1111111110000...00 = -\infty$
- $X11111111[\text{non-zero}] = \text{NaN}$
 - e.g., square root of a negative number
- $X0000000000000...00 = 0$
 - ...there's actually a positive zero and a negative zero

Floating-Point Addition

- Consider a 4-digit decimal example
 - $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
 - Shift number with smaller exponent
 - $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
 - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
 - 1.0015×10^2
- 4. Round and renormalize if necessary
 - 1.002×10^2

Floating-Point Addition

- Now consider a 4-digit binary example
 - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$ (0.5 + -0.4375)
- 1. Align binary points
 - Shift number with smaller exponent
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- 2. Add significands
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
 - $1.000_2 \times 2^{-4}$, with no over/underflow
- 4. Round and renormalize if necessary
 - $1.000_2 \times 2^{-4}$ (no change) = 0.0625

FP Instructions in MIPS

- FP hardware is coprocessor 1
 - Adjunct processor that extends the ISA
- Separate FP registers
 - 32 single-precision: \$f0, \$f1, ... \$f31
 - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
 - Release 2 of MIPS ISA supports 32×64 -bit FP reg's
- FP instructions operate only on FP registers
 - Programs generally don't do integer ops on FP data, or vice versa
 - More registers with minimal code-size impact
- FP load and store instructions
 - lwc1, ldc1, swc1, sdc1
 - e.g., ldc1 \$f8, 32(\$sp)

FP Instructions in MIPS

- Single-precision arithmetic
 - `add.s`, `sub.s`, `mul.s`, `div.s`
 - e.g., `add.s $f0, $f1, $f6`
- Double-precision arithmetic
 - `add.d`, `sub.d`, `mul.d`, `div.d`
 - e.g., `mul.d $f4, $f4, $f6`
- Single- and double-precision comparison
 - `c.xx.s`, `c.xx.d` (*xx* is `eq`, `lt`, `le`, ...)
 - Sets or clears FP condition-code bit
 - e.g. `c.lt.s $f3, $f4`
- Branch on FP condition code true or false
 - `bc1t`, `bc1f`
 - e.g., `bc1t TargetLabel`

FP Example: °F to °C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

- fahr in \$f12, result in \$f0, literals in global memory space

- Compiled MIPS code:

```
f2c: lwc1    $f16, const5($gp)  
     lwc1    $f18, const9($gp)  
     div.s   $f16, $f16, $f18  
     lwc1    $f18, const32($gp)  
     sub.s   $f18, $f12, $f18  
     mul.s   $f0, $f16, $f18  
     jr      $ra
```

FP Example: Array Multiplication

- $X = X + Y \times Z$ (initialize X to all zeros...)
 - All 32×32 matrices, 64-bit double-precision elements

- C code:

```
void mm (double x[][],
         double y[][], double z[][]) {
    int i, j, k;
    for (i = 0; i != 32; i = i + 1)
        for (j = 0; j != 32; j = j + 1)
            for (k = 0; k != 32; k = k + 1)
                x[i][j] = x[i][j]
                    + y[i][k] * z[k][j];
}
```

- Addresses of x, y, z in $\$a0, \$a1, \$a2$, and
 i, j, k in $\$s0, \$s1, \$s2$

FP Example: Array Multiplication

■ MIPS code:

| | | | |
|-----|------|------------------|------------------------------------|
| | li | \$t1, 32 | # \$t1 = 32 (row size/loop end) |
| | li | \$s0, 0 | # i = 0; initialize 1st for loop |
| L1: | li | \$s1, 0 | # j = 0; restart 2nd for loop |
| L2: | li | \$s2, 0 | # k = 0; restart 3rd for loop |
| | sll | \$t2, \$s0, 5 | # \$t2 = i * 32 (size of row of x) |
| | addu | \$t2, \$t2, \$s1 | # \$t2 = i * size(row) + j |
| | sll | \$t2, \$t2, 3 | # \$t2 = byte offset of [i][j] |
| | addu | \$t2, \$a0, \$t2 | # \$t2 = byte address of x[i][j] |
| | l.d | \$f4, 0(\$t2) | # \$f4 = 8 bytes of x[i][j] |
| L3: | sll | \$t0, \$s2, 5 | # \$t0 = k * 32 (size of row of z) |
| | addu | \$t0, \$t0, \$s1 | # \$t0 = k * size(row) + j |
| | sll | \$t0, \$t0, 3 | # \$t0 = byte offset of [k][j] |
| | addu | \$t0, \$a2, \$t0 | # \$t0 = byte address of z[k][j] |
| | l.d | \$f16, 0(\$t0) | # \$f16 = 8 bytes of z[k][j] |

...

FP Example: Array Multiplication

...

| | | |
|-------|---------------------|------------------------------------|
| sll | \$t0, \$s0, 5 | # \$t0 = i*32 (size of row of y) |
| addu | \$t0, \$t0, \$s2 | # \$t0 = i*size(row) + k |
| sll | \$t0, \$t0, 3 | # \$t0 = byte offset of [i][k] |
| addu | \$t0, \$a1, \$t0 | # \$t0 = byte address of y[i][k] |
| l.d | \$f18, 0(\$t0) | # \$f18 = 8 bytes of y[i][k] |
| mul.d | \$f16, \$f18, \$f16 | # \$f16 = y[i][k] * z[k][j] |
| add.d | \$f4, \$f4, \$f16 | # \$f4 = x[i][j] + y[i][k]*z[k][j] |
| addiu | \$s2, \$s2, 1 | # k = k + 1 |
| bne | \$s2, \$t1, L3 | # if (k != 32) go to L3 |
| s.d | \$f4, 0(\$t2) | # x[i][j] = \$f4 |
| addiu | \$s1, \$s1, 1 | # j = j + 1 |
| bne | \$s1, \$t1, L2 | # if (j != 32) go to L2 |
| addiu | \$s0, \$s0, 1 | # i = i + 1 |
| bne | \$s0, \$t1, L1 | # if (i != 32) go to L1 |

Accurate Arithmetic

- IEEE Std 754 specifies additional rounding control
 - Extra bits of precision (guard, round, sticky)
 - Choice of rounding modes
 - Allows programmer to fine-tune numerical behavior of a computation
- Not all FP units implement all options
 - Most programming languages and FP libraries just use defaults
- Trade-off between hardware complexity, performance, and market requirements

Associativity

- Parallel programs may interleave operations in unexpected orders
 - Assumptions of associativity may fail

| | | $(x+y)+z$ | $x+(y+z)$ |
|---|-----------|-----------|-----------|
| x | -1.50E+38 | 0.00E+00 | -1.50E+38 |
| y | 1.50E+38 | | 1.50E+38 |
| z | 1.0 | 1.0 | |
| | | 1.00E+00 | 0.00E+00 |

- Need to validate parallel programs under varying degrees of parallelism

Who Cares About FP Accuracy?

- Important for scientific code
 - But for everyday consumer use?
 - "My bank balance is out by 0.0002¢!" ☹
- The Intel Pentium FDIV bug
 - The market expects accuracy
 - See Colwell, *The Pentium Chronicles*